

OSA CON 25



Micromegas

Unified observability for video games

Marc-Antoine Desroches
Technical Architect
madesroches@gmail.com

November 4-5, 2025



Disclaimer

These are my personal opinions and experiences.
I'm not speaking for my employer.



The Objective

Quantify **how often** and **how bad** issues are,
with enough context to **fix** them
no need to reproduce.



The Problem

Traditional tooling force a choice

🔍 High-frequency debugging tools

Great detail, but you need to reproduce the bug yourself

📊 Low-frequency analytics tools

Will report statistics, not detailed traces

We refuse to choose.



The Challenge

- Video games at 60fps generate enormous telemetry volumes:
- ~10s of log entries/sec
- ~1k measures/sec - 10s/frame
- ~60k-200k CPU trace events/sec - 1000s/frame
- Thousands of concurrent processes
- Unified system for logs, metrics, and traces



Operating Costs - real production example

Data volume

- Retention: 90 days
- 9 billion log entries
- 275 billion metrics
- 165 billion sampled trace events

- Total 449 billion events
- Spikes of 266M events/minute

Monthly Cost breakdown

- Compute ~\$300
- PostgreSQL ~\$200
- S3 Storage: ~\$500

Total ~\$1,000/month




How

Recording & querying high volumes of data can be cheap


- Stage 1. Low-overhead always-on instrumentation - 20ns/event
- Stage 2. Cheap ingestion & storage
- Stage 3. Daemon processes low-frequency streams (logs & metrics)
- Stage 4. Analytics service executes SQL queries
- Stage 5. Frugal user interface

Nothing in here is specific to games

Architecture Overview - Ingestion



Architecture Overview - Analytics





Stage 1: Low-Overhead Instrumentation

Optimized instrumentation libraries

How?

- Events are **tiny** (a few bytes)
- Events can contain **pointers** to avoid repetition (metric names, file paths, ...)
- Instrumented threads **only queue events** - work done in background thread
 - overhead ~**20 ns** / cpu trace event in calling thread
- **Sampling logic applies to event batches**
 - ex: does this section of the trace include slow frames?
- **Serialization is mostly memcpy**
 - Event layout is the **native memory format**
 - Event layout changes from one process to another
- Fast compression using **LZ4**



Code Example: Rust Instrumentation

```
use micromegas_tracing::prelude::*;

#[span_fn]
async fn process_request(user_id: u32) -> Result<Response> {
    info!("request user_id={user_id}");
    let begin_ticks = now();
    let response = handle_request(user_id).await?;
    let end_ticks = now();
    let duration = end_ticks - begin_ticks;
    imetric!("request_duration", "ticks", duration as u64);
    info!("response status={}", response.status());
    Ok(response)
}
```



Code Example: Unreal Instrumentation

```
#include "MicromegasTracing/Macros.h"

float AMyActor::TakeDamage(float Damage, ...)
{
    MICROMEGAS_SPAN_FUNCTION("Combat");
    float ActualDamage = Super::TakeDamage(...);
    MICROMEGAS_FMETRIC("Combat",
        MicromegasTracing::Verbosity::High,
        TEXT("DamageDealt"), TEXT("points"),
        ActualDamage);
    return ActualDamage;
}
```



Stage 2: Ingestion & Storage

Simple, horizontally scalable design

HTTP service accepts LZ4-compressed payloads

- Metadata → PostgreSQL (for fast lookups)
- Payloads → S3 (for cheap storage)

Data lake: optimized for cheap writes



Stages 3 & 4: daemon & analytics

Bridge between **data lake** (cheap writes) and **lakehouse** (fast reads)

- Extract binary payload data in compressed **custom format**
- Transform into Parquet
- Let DataFusion execute SQL queries on the parquet files



Stage 3: Daemon

Different streams, different strategies

- Logs (low frequency)
 - Process eagerly → Parquet
- Metrics (medium frequency)
 - Process eagerly → Parquet
- Keeps aggregated views up to date
- Ignores CPU traces (very high frequency)



Incremental Data Reduction

Example: SQL-defined log_stats view

```
SELECT date_bin('1 minute', time) as time_bin,  
       process_id,  
       level,  
       target,  
       count(*) as count  
  FROM log_entries  
 WHERE insert_time >= '{begin}'  
   AND insert_time < '{end}'  
 GROUP BY process_id, level, target, time_bin
```

Allows to query data over multiple days



Stages 4 : Analytics

Serves SQL queries over Arrow FlightSQL

- Transform payloads into Parquet if daemon did not already do it
 - CPU traces
 - Process-specific views
- Let DataFusion SQL engine do its magic
- DataFusion is augmented by extensions
 - JSONB
 - Histograms (median, p99, ...)




Stage 5: User Interfaces

Three main interfaces

- Grafana
 - dashboards and alerting
- Jupyter Notebooks
 - Python API for data exploration
- Perfetto
 - Deep trace visualization



Grafana Dashboard





Jupyter Notebooks

Python API
for flexible
data
exploration
and custom
analysis

A screenshot of a Jupyter Notebook cell. The code imports datetime, pandas, and micromegas, sets a display option, connects to a client, and runs a query to select log entries from the last 7 days. The output shows the CPU times and a table of log entries with columns: index, time, level, exe, target, and msg. The table lists 10 rows of log data.

```
: import datetime
import pandas as pd
import micromegas
pd.set_option('display.max_rows', 30)
client = micromegas.connect()

: %time
now = datetime.datetime.now(datetime.timezone.utc)
begin = now - datetime.timedelta(days=7)
end = now
sql = """
SELECT time, level, exe, target, msg
FROM log_entries
WHERE level <= 4
ORDER BY time DESC
limit 10
"""

client.query(sql, begin, end)


CPU times: user 0 ns, sys: 2.35 ms, total: 2.35 ms
Wall time: 62.7 ms
```

	time	level	exe	target	msg
0	2025-10-31 11:24:33.573681118+00:00	4	/home/mad/micromegas/rust/target/debug/telemet...	micromegas_analytics::lakehouse::write_partition	[PARTITION_WRITE_COMMIT] view=streams/global t...
1	2025-10-31 11:24:33.571009166+00:00	4	/home/mad/micromegas/rust/target/debug/telemet...	micromegas_analytics::response_writer	[PARTITION_WRITE_START] view=streams/global ti...
2	2025-10-31 11:24:33.570000195+00:00	4	/home/mad/micromegas/rust/target/debug/telemet...	micromegas_analytics::response_writer	creating empty partition record for [streams, ...
3	2025-10-31 11:24:33.563986492+00:00	4	/home/mad/micromegas/rust/target/debug/telemet...	micromegas_analytics::response_writer	writing [streams, global] 2025-10-31T11:24:31+...
4	2025-10-31 11:24:33.563922756+00:00	4	/home/mad/micromegas/rust/target/debug/telemet...	micromegas_analytics::response_writer	[2025-10-31T11:24:31+00:00, 2025-10-31T11:24:3...
5	2025-10-31 11:24:33.557295458+00:00	4	/home/mad/micromegas/rust/target/debug/telemet...	micromegas_analytics::lakehouse::write_partition	[PARTITION_WRITE_COMMIT] view=processes/global...
6	2025-10-31 11:24:33.554438214+00:00	4	/home/mad/micromegas/rust/target/debug/telemet...	micromegas_analytics::response_writer	[PARTITION_WRITE_START] view=processes/global ...
7	2025-10-31 11:24:33.553383338+00:00	4	/home/mad/micromegas/rust/target/debug/telemet...	micromegas_analytics::response_writer	creating empty partition record for [processes...
8	2025-10-31 11:24:33.544960836+00:00	4	/home/mad/micromegas/rust/target/debug/telemet...	micromegas_analytics::response_writer	writing [processes, global] 2025-10-31T11:24:3...
9	2025-10-31 11:24:33.544897840+00:00	4	/home/mad/micromegas/rust/target/debug/telemet...	micromegas_analytics::response_writer	[2025-10-31T11:24:31+00:00, 2025-10-31T11:24:3...



Perfetto Trace Viewer

Detailed CPU trace analysis





Thank You

Micromegas would not be possible without open source

- FlightSQL, DataFusion, Apache Arrow, Parquet
- InfluxData
- PostgreSQL
- Rust

And many other amazing projects



Micromegas is Open Source

★ <https://github.com/madesroches/micromegas>

- Drop a star (always makes my day!)
- Try it out, use it as a library, copy the code
- Open an issue, tell me what's missing
- Share your use cases
- madesroches@gmail.com



Q&A